



Rootless containers from scratch

Liz Rice

VP Open Source Engineering, Aqua Security

@lizrice

What is a container?

What is a rootless container?

By default, containers run as root

The same root as on the host



✖ **Warning:**

Adding a user to the “docker” group grants them the ability to run containers which can be used to obtain root privileges on the Docker host. Refer to [Docker Daemon Attack Surface](#) for more information.



✓ **Note:**

To install Docker without root privileges, see [Run the Docker daemon as a non-root user \(Rootless mode\)](#).

Rootless mode is currently available as an experimental feature.

Namespaces

Limit what a process can **see**

- Unix Timesharing System
- Process IDs
- Mounts
- Network
- InterProcess Comms
- User IDs

Namespaces

Limit what a process can **see**

- Unix Timesharing System
- Process IDs
- Mounts
- Network
- InterProcess Comms
- User IDs

man user_namespaces

Starting in Linux 3.8, **unprivileged processes can create user namespaces**

man user_namespaces

Starting in Linux 3.8, unprivileged processes can create user namespaces

The child process created ... with the CLONE_NEWUSER flag starts out with a **complete set of capabilities** in the new user namespace.

man user_namespaces

Starting in Linux 3.8, unprivileged processes can create user namespaces

The child process created ... with the CLONE_NEWUSER flag starts out with a complete set of capabilities in the new user namespace.

If CLONE_NEWUSER is specified along with other CLONE_NEW* flags ... the user namespace is guaranteed to be created first, giving the child ... **privileges over the remaining namespaces** created by the call.

Host User IDs

0

1

2

...

1000

1001

1002

1003

...

Namespace User IDs

0

1

2

...

size

```

func main() {
    switch os.Args[1] {
        case "run":
            run()
        case "child":
            child()
        default:
            panic("Missing argument 1")
    }
}

func run() {
    fmt.Printf("Running %v as user %d in process %d\n", os.Args[2:], os.Geteuid(), os.Getpid())

    cmd := exec.Command("/proc/self/exe", append([]string{"child"}, os.Args[2:]...))
    cmd.Stdout = os.Stdout
    cmd.Stderr = os.Stderr
    cmd.Stdin = os.Stdin
    cmd.SysProcAttr = &syscall.SysProcAttr{
        Cloneflags: syscall.CLONE_NEWUTS | syscall.CLONE_NEWUSER | syscall.CLONE_NEWNS | syscall.CLONE_NEWPID,
        UidMappings: []syscall.SysProcIDMap{
            {
                ContainerID: 0,
                HostID:      1000,
                Size:         1}},
    }
    must(cmd.Run())
}

```

```

func child() {
    fmt.Printf("Running %v as user %d in process %d\n", os.Args[2:], os.Geteuid(), os.Getpid())
    fmt.Printf("Capabilities: %s\n", showCaps())
    must(Chterm("/home/vagrant/alpinefs"))
    must(Cbdir("/"))
    must(Mount("proc", "/proc", "proc", 0, ""))
}

```



```
func child() {
    fmt.Printf("Running %v as user %d in process %d\n", os.Args[2:], os.Geteuid(), os.Getpid())

    must(syscall.Chroot("/home/vagrant/alpinefs"))
    must(os.Chdir("/"))
    must(syscall.Mount("proc", "proc", "proc", 0, ""))

    cmd := exec.Command(os.Args[2], os.Args[3:]...)
    cmd.Stdout = os.Stdout
    cmd.Stderr = os.Stderr
    cmd.Stdin = os.Stdin
    must(cmd.Run())

    must(syscall.Unmount("proc", 0))
}

func must(err error) {
    if err != nil {
        panic(err)
    }
}
```

● Rootless container support

Changelog

For official release notes for Docker Engine CE, visit the [release notes page](#).

20.10.0

Rootless

- rootless: graduate from experimental [moby/moby#40759](#)


● Rootless container support

 [kubernetes](#) / [enhancements](#)

[Code](#) [Issues 213](#) [Pull requests 116](#) [Actions](#) [Projects 2](#) [Security](#) [Insights](#)

keps/127: Support User Namespaces #2101

[Open](#)

mauriciovasque... wants to merge 4 commits into [kubernetes:master](#) from [kinvolk:mauricio/usersns_proposal_upstream](#) 

Thank you

github.com/rootless-containers/rootlesskit

github.com/lizrice/containers-from-scratch

